



Runtime Security for Agentic Systems A Practical Two-Plane Approach for OpenClaw-Class Agents

Frank Lyonnet, Antoine Clerget, Kave Salamatian

► To cite this version:

Frank Lyonnet, Antoine Clerget, Kave Salamatian. Runtime Security for Agentic Systems A Practical Two-Plane Approach for OpenClaw-Class Agents. 2026. ⟨hal-05607320⟩

HAL Id: hal-05607320

<https://hal.science/hal-05607320v1>

Preprint submitted on 29 Apr 2026

HAL is a multi-disciplinary open access archive for the deposit and dissemination of scientific research documents, whether they are published or not. The documents may come from teaching and research institutions in France or abroad, or from public or private research centers.

L'archive ouverte pluridisciplinaire **HAL**, est destinée au dépôt et à la diffusion de documents scientifiques de niveau recherche, publiés ou non, émanant des établissements d'enseignement et de recherche français ou étrangers, des laboratoires publics ou privés.



ETALAB - Open licence

Runtime Security for Agentic Systems

A Practical Two-Plane Approach for OpenClaw-Class Agents

Authors: Frank Lyonnet, Antoine Clerget, Kave Salamatian

Status: Working draft

Version: 2026-03-17

Abstract

Tool-using AI agents expose a runtime security gap: pre-install scanning and configuration hardening are necessary, but they cannot fully characterize behavior that emerges during execution (e.g., indirect prompt injection and tool poisoning). To reduce that residual gap, we present a two-plane security model that correlates reasoning-plane observations (independently read session transcripts) with system-plane telemetry to detect intent-behavior divergence.

The contribution is architectural rather than a claim of universal prevention. EDAMAME accepts agent-tagged behavioral-model slices from multiple reasoning-plane producers, merges them into one observer-owned correlation window, and evaluates those slices against process-attributed host telemetry and model-independent guardrails. The same control loop supports explicit behavioral slices (`upsert_behavioral_model`) and raw-session ingest (`upsert_behavioral_model_from_raw_sessions`) for workstation paths such as Cursor and Claude Code.

We present a quantitative result that is narrower than a full end-to-end agent evaluation. These snapshot results support the feasibility of the correlation loop as an early-warning mechanism. BadAgentUse uses explicit behavioral-model slice injection to exercise EDAMAME’s downstream correlation loop under accelerated benchmark cadence, while end-to-end transcript-publication paths are validated separately. On the current reproducible branch snapshot (468 valid runs across 10 seeds, with skipped runs tracked separately in the readiness scorecard), precision is 1.000 (CI95: [0.988, 1.000]) and recall is 0.911 (CI95: [0.876, 0.937]). Median detection latency is 51.6 s (p95: 66.9 s). We keep three limitations explicit: reactive latency; residual adversarial-evasion misses dominated by explicitly versioned `DESIGNED_MISS` living-off-the-land scenarios that remain in the recall denominator, plus one remaining non-designed miss in network-intent-correlation; and living-off-the-land behavior that stays within allowed process/destination pairs.

Keywords

agent security, runtime monitoring, prompt injection, tool poisoning, Model Context Protocol (MCP), intent-behavior correlation, telemetry, reproducibility

1. Introduction

Agent systems now perform high-impact operations through tools and MCP integrations: network access, file mutation, credential-adjacent actions, and persistent memory updates. Risk therefore depends on what the agent actually does at runtime, not only on what is statically configured before runtime.

A recent publicly discussed incident involving a Meta AI security researcher and an OpenClaw deployment that reportedly executed unintended mailbox-deletion actions (Bort, 2026) illustrates the practical problem: loss of control can happen during normal runtime, even in technically sophisticated environments. Our model targets that runtime-loss-of-control phase by providing independent host-side observation and divergence alerts. However, it does not replace authorization design or pre-action approval gates; if broad destructive permissions are already granted, prevention still depends on those controls.

We characterize the problem in terms of three **control strata** to avoid terminology collision with our model’s two planes:

1. **S1: Supply-chain controls** (e.g., signed artifacts, provenance/SBOM checks).
2. **S2: Configuration controls** (e.g., localhost MCP bind + PSK auth, role and tool scoping, sandboxing, policy hardening).
3. **S3: Runtime behavior controls** (e.g., a declared read-only task followed by undeclared egress, detected via intent-behavior consistency checks).

The residual runtime-only failures that remain after S1/S2 controls are what we call the **S3 gap**.

Together, S1/S2 can stop many failures. For example, an internet-exposed OpenClaw gateway (SecurityScorecard STRIKE, 2026) is often prevented by correct bind/auth configuration, which is fundamentally a configuration-control problem. However, S1/S2 do not fully address runtime divergence: in CVE-2026-25253, a browser token exfiltration path can execute during normal user interaction even when baseline setup appears acceptable. Similar runtime-only gaps arise in coding agents operating over production repositories, where unauthorized access to git-managed secrets, SSH keys, or CI tokens can produce outsized impact.

This paper addresses the S3 gap with a two-plane model: compare what the reasoning plane declares with what the system plane observes on the host (defined in Section 3). The approach highlights a structural opportunity for security tooling vendors (here, EDAMAME Posture) and AI platform providers (here, OpenClaw, Cursor, and Claude Code) to collaborate through MCP: the security tool provides system-plane ground truth that no AI framework includes natively, while the AI platform provides the runtime observation surface (session transcripts, memory, tool invocation history) that no external monitor can access alone. Neither plane is sufficient in isolation; the two-plane correlation is

the necessary security primitive.

The core model claim is two-plane correlation. In the reference instantiation, the system plane decomposes into three adjacent runtime loops (Section 3). The first – the **divergence / intent loop** – is the paper’s primary contribution: it correlates a merged behavioral model against process-attributed telemetry and flags observations not explained by declared intent. The remaining two loops address capabilities that are closer to traditional EDR (endpoint detection and response) than to intent correlation: a **vulnerability / safety-floor loop** that evaluates model-independent guardrail findings (for example `/tmp` lineage, credential-file-plus-egress, and exposed-listener detectors) even when no behavioral model exists, and an **advisor / remediation loop** that groups posture, breach, and LAN findings into operator-facing todos and reversible actions. These two loops complement the intent-correlation contribution but are not themselves novel; their inclusion ensures that the system plane provides useful security coverage even before a behavioral model is available.

By design, this first-generation rule set has a blind spot for living-off-the-land behavior when both process and destination are already authorized; we surface that limitation explicitly in Sections 7 and 9.

Throughout the paper, we distinguish between what is implemented in the reference runtime, what is validated by integration tests, and what is measured in the current quantitative snapshot. That distinction matters because the current benchmark isolates the downstream correlation loop rather than the full transcript-to-model publication path.

1.1 Contributions (Novelty)

We separate contributions into a model part and an implementation/evaluation part.

Model contributions (Part I):

- **Two-plane runtime correlation model.** A control loop for correlating observed agent behavior (reasoning plane, via session transcript observation) with independent host telemetry (system plane).
- **Model-level requirements.** Process-attributed telemetry, observer-owned state, and a signal hierarchy are presented as requirements of the model, not as product-specific implementation details.
- **MCP-aware security framing.** We use MCP as a typed boundary between reasoning-plane publishers and an external security observer, not only as a transport convenience.

Implementation and evaluation contributions (Part II):

- **EDAMAME instantiation with two producer modes.** A concrete deployment path over local, PSK-authenticated MCP with rollback-capable action handling and merged multi-agent reasoning-plane inputs.

The same runtime accepts explicit prebuilt behavioral windows and raw-session inputs converted inside EDAMAME.

- **Trace-backed benchmark protocol.** Every verdict is tied to on-disk telemetry artifacts for auditability.
- **Current empirical snapshot of the downstream correlation loop.** The quantitative section measures explicit-slice benchmark mode on a live 50-scenario suite, including explicit designed-miss scenarios, while keeping end-to-end producer validation separate.

Part I — Two-Plane Model

Part I is organized as follows: Section 2 positions this work in prior runtime security literature; Section 3 defines the two-plane model; Section 4 states threat assumptions and boundaries; and Section 5 gives concise properties of the decision rules as a bridge to implementation. Readers primarily interested in model mechanics and applications may skip Section 2 and proceed directly to Section 3.

2. Background and Related Work

2.1 Attack Surface

Indirect prompt injection blurs data and instructions in LLM-integrated applications, enabling remote compromise via retrieved content (Greshake et al., 2023). Benchmarks such as InjecAgent show tool-integrated agents remain vulnerable to indirect prompt injection (Zhan et al., 2024).

MCP expands the attack surface by standardizing tool descriptors and invocations. Attacks embedded in tool metadata (tool poisoning, descriptor mutation) achieve attack success rates up to 72.8% on capable models (MCPTox; Luo et al., 2025). Mitigations include descriptor signing, semantic vetting (Jamshidi et al., 2025), and multi-stage neural detection pipelines achieving 96% accuracy on adversarial prompts (MCP-Guard; Chen et al., 2025). MCP’s own specification documents authorization and security best practices, including confused deputy risks in proxying designs (MCP Spec, 2025). MCP Security Bench provides a 2,000-instance, 12-attack-type evaluation suite for systematically measuring tool-level defenses (Wang et al., 2025).

Goal drift – an agent’s tendency to deviate from original objectives over extended execution – is now measurable in controlled evaluations. Zhou et al. (2025) show that all evaluated models exhibit some drift under autonomous operation, even when strong scaffolding substantially reduces it.

2.2 Runtime Defense Landscape

The 2025-2026 research cycle has produced a rich landscape of runtime defense systems. We organize them by their primary signal source.

Reasoning-plane defenses monitor the agent’s internal state (prompts, chain of thought, tool call sequences) to detect compromise:

- AgentArmor (Gao et al., 2025) treats runtime traces as program objects, constructs CFG/DFG/PDG views, and applies type-checking to trust-boundary flows. It reports 97% defense success (3% attack success) on AgentDojo with modest utility impact.
- LlamaFirewall (Meta, 2025) provides three production guardrails: PromptGuard 2 (BERT-based jailbreak detection), Agent Alignment Checks (chain-of-thought auditing for prompt injection and goal misalignment), and CodeShield (online static analysis). It is deployed in production at Meta.
- AegisLLM (Li et al., 2025) uses cooperative multi-agent defense where orchestrator, deflector, responder, and evaluator agents collaborate for safe outputs, reporting a 51% improvement on jailbreak benchmarks without retraining.
- A2AS (Wallarm et al., 2025) proposes the BASIC security model: behavior certificates, authenticated prompts, security boundaries, in-context defenses, and codified policies. The initiative has broad industry sponsorship and is emerging as a framework-level governance model.

Hybrid defenses combine reasoning-state signals and runtime action signals:

- MI9 (Singh et al., 2025) is the first integrated runtime governance framework, combining goal-conditioned drift detection, FSM-based conformance engines, continuous authorization monitoring, and graduated containment. Reports 99.81% detection on synthetic scenarios.
- AgentSentinel (Hu et al., 2025) intercepts all sensitive operations within agent services, halting execution until a security audit correlating task context with system traces completes. Achieves 79.6% defense success across 60 attack scenarios in 6 categories.
- Pro2Guard (Zhang et al., 2025) learns Discrete-Time Markov Chains (DTMCs) from execution traces and proactively predicts unsafe states before they occur, with PAC-style correctness bounds. Reported defense success is 97.4% (2.6% unsafe outcomes).

Formal verification approaches provide provable safety properties:

- ShieldAgent (Xu et al., 2025) constructs safety policies as verifiable rule circuits and generates shielding plans with executable code for formal verification. Achieves 90.1% recall on safety violations.
- Agent-C (Kamath et al., 2025) enforces temporal safety constraints via SMT solving and constrained generation, achieving 100% conformance on leading-model evaluations.
- VET (Li et al., 2025) establishes verifiable execution traces via trusted hardware, cryptographic proofs, and notarized TLS transcripts with practical overhead under 3x for API calls.
- AgentGuard (Koohestani, 2025) provides runtime verification via proba-

bilistic model checking and Dynamic Probabilistic Assurance for quantitative agent properties.

Chain-of-thought monitoring research reveals important limitations:

- Models can generate benign-sounding CoT explanations that mask misaligned intentions, causing monitors to fail (Thought Crime; Rogers et al., 2025).
- Agent awareness of monitoring is the single strongest evasion advantage – detection rates drop sharply when agents know they are monitored (Baker et al., 2025).
- Hybrid protocols scoring both reasoning and actions consistently surpass either alone, with 2x higher detection for subtle deception scenarios (CoT Red-Handed; Goldowsky-Dill et al., 2025).

2.3 Positioning

Table 2 summarizes the landscape. Our system combines MCP-native deployment, multiplatform process-attributed telemetry, and intent correlation. The closest comparator is AgentSentinel, which also correlates task context with OS-level traces via eBPF/LSM. Key differentiators:

- **Multiplatform observer:** our observer (EDAMAME Posture) works across Linux, macOS, and Windows (Section 4.3), while AgentSentinel requires Linux eBPF/LSM. eBPF enhances attribution resolution on Linux but is not required.
- **MCP-native deployment surface:** telemetry is exposed as standard MCP tools over Streamable HTTP with PSK authentication, usable by any MCP-compatible agent without custom integration.
- **Observer-owned merge semantics:** EDAMAME merges multiple agent-tagged contributor slices centrally, so no single reasoning-plane producer owns the correlation window or verdict history.
- **Correlation vs. interception:** AgentSentinel halts execution at sensitive operations for synchronous audit; our system primarily correlates observed behavior after it appears in telemetry, with an optional proactive layer for MCP-gated actions. This trades immediate blocking for lower runtime friction and broader platform compatibility.
- **Measurement scope (explicit).** The 51.6 s median detection delay reported in this draft comes from the current accelerated branch snapshot, not from a steady-state production default cadence.

Its primary limitations relative to the field are partial reasoning-plane coverage (transcript-only, no CoT introspection; Section 9) and absence of formal safety proofs.

System	Plane	MCP	Platforms	Eval. Scale
AgentArmor	Reasoning	No	Any	AgentDojo (97%)
LlamaFirewall	Reasoning	No	Any	Meta internal
AegisLLM	Reasoning	No	Any	Jailbreak bench.
A2AS BASIC	Reasoning	Yes	Any	Framework-level
MI9	Hybrid	No	Any	Synthetic (99.81%)
AgentSentinel	Hybrid	No	Linux (eBPF)	60 scen., 6 cat.
Pro2Guard	Hybrid	No	Any	Household + AV
ShieldAgent	Formal	No	Any	Web environments
Agent-C	Formal	No	Any	Multiple bench.
VET	Trace	No	Any	Black-box API
MCP-Guard	Metadata	Yes	Any	70,448 samples
This work	Hybrid	Yes	Lin/Mac/Win	50 + 2 integ.

A2AS BASIC is conceptual at the MCP-native layer. **This work** is hybrid because reasoning-plane sessions and system-plane telemetry are observed independently and merged only inside the observer-owned runtime. In this table, **50 scen. + 2 integ.** means 50 benchmark scenarios across six benchmark categories plus two integration-only signal families (posture and LAN) that are validated outside the confusion matrix.

Formal properties reported by prior systems: - A2AS (behavioral certificates) - MI9 (FSM conformance) - Pro2Guard (PAC-correctness) - ShieldAgent (rule circuits) - Agent-C (SMT temporal constraints) - VET (cryptographic proof) - MCP-Guard (neural detector, reported 96%)

Deployed systems in this survey scope: - LlamaFirewall (Meta production) - A2AS (emerging standard) - This work (research deployment in Lima VM) - All others are currently research-only artifacts.




Figure 1: Agent Runtime Security Timeline

We do not claim to be the first system to expose OS- or eBPF-derived telemetry through MCP for analysis. Public observability projects such as the Inspektor Gadget MCP Server and AKS-MCP observability tooling already show that telemetry-via-MCP is practical. Our narrower claim is that this work uses MCP-exposed host telemetry to produce security intent-behavior divergence verdicts for agent runtimes under an observer-owned merge contract, and that the current draft evaluates that correlation loop on trace-backed scenarios.¹ The trade-off is that system-plane monitoring remains reactive (Section 3.6) and reasoning-plane coverage remains transcript-only, without CoT introspection (Section 9).

2.4 Reflection as Relevant Prior Art

The model is also related to reflection-oriented programming, where a program observes and reasons about its own behavior through a meta-level interface. Sobel and Friedman (1996) frame reflection as separating base-level execution from meta-level supervision. Our two-plane model adopts a similar separation: the reasoning plane executes and declares intent; the system plane independently observes effects and can challenge that declaration. We do not claim formal equivalence to classic reflection systems, but the analogy is useful for conceptual grounding.

3. Two-Plane Architecture (Reasoning Plane and System Plane)

- **Reasoning plane:** what the agent declares it is doing and why.
- **System plane:** what the host actually does.

At a high level, the control loop is:

```
declare intent ->
derive expected effects ->
observe host behavior ->
correlate ->
compute deterministic evidence ->
optionally adjudicate soft signals with an LLM ->
reconcile critical-signal guardrails ->
emit final divergence verdict
```

The architecture maps reasoning-plane intent declarations to system-plane telemetry correlation, and emits divergence when observed behavior is not explained by declared intent. In the deployed system, the behavioral model consumed by the divergence engine can be a merged window assembled from

¹Survey scope (as of February 2026): arXiv, ACM DL, IEEE Xplore, and USENIX proceedings through February 2026 using keywords including “agent runtime monitoring”, “MCP security”, “tool-use telemetry”, and “eBPF agent”, plus manual review of MCP specifications and ecosystem tooling.

Two-Plane Architecture with Three Internal Runtime Loops




Figure 2: Two-Plane Architecture with Reasoning-Plane Declarations and System-Plane Runtime Loops

multiple agent-tagged contributors (for example OpenClaw plus Cursor and Claude Code) rather than a single producer. Producer modes are additive: OpenClaw commonly uses the prebuilt-window publication path, while workstation integrations such as Cursor and Claude Code can forward raw reasoning sessions so EDAMAME generates the contributor slice internally before merge.

The operational cadence is shared-timer-driven: the extrapolator cron publishes behavioral models on a fixed interval, while EDAMAME evaluates three adjacent runtime loops internally and handles divergence alerting through its own notification subsystem. The timer structure is important to portability: it avoids dependence on framework-specific synchronous hooks and keeps behavior consistent across deployments. We considered two alternatives – interception-first synchronous gating (lower delay but higher runtime friction and tighter platform coupling) and a single merged polling loop (simpler but weaker separation between model publication and engine evaluation). The timer-plus-internal-notification design gave the best operator clarity while preserving framework-agnostic deployment.

Cadence is an implementation parameter, not a model contribution. In the current reference runtime, production defaults use 300 s divergence and vulnerability intervals plus a $\ast/5$ extrapolator cron cadence. Executive demo scripts accelerate the internal loops to 60 s and the extrapolator cron to $\ast/2$, while the explicit-slice benchmark helper can accelerate the engine further because it injects the behavioral slice directly before polling for verdicts. This separation is what reconciles the production default timer story with the sub-minute latency numbers reported in Section 7.

The paper’s model claim remains two-plane correlation. The three-loop split is an implementation decomposition of the system plane in the reference instantiation:

1. **Divergence / intent loop.** Correlates the merged behavioral model against raw process-attributed telemetry and emits intent-relative evidence when observations are not explained by declared behavior. This is the loop that realizes the two-plane model contribution.
2. **Vulnerability / safety-floor loop.** Evaluates model-independent guardrail findings and runtime incident checks (for example `/tmp` lineage, credential-file-plus-egress, exposed listeners, and other vulnerability detectors), including cases where the behavioral model is missing or stale.
3. **Advisor / remediation loop.** Groups posture, breach, LAN, and high-confidence telemetry findings into operator-facing todos and reversible actions.

Loops 2 and 3 are closer to traditional endpoint-detection-and-response (EDR) capabilities than to intent correlation; they are included so the system plane provides useful security coverage even before a behavioral model is available, but they are not the paper’s novel contribution.

This separation is deliberate: telemetry confidence labels such as `anomalous`

or `blacklisted` are useful operational signals, but they are not treated as intent-relative divergence evidence by themselves. Scenarios are judged from engine verdicts read via `get_divergence_verdict`; the final operator verdict may include optional LLM adjudication for soft-signal cases, but deterministic safety-floor and vulnerability-loop findings remain non-downgradable.




Figure 3: Shared-Timer Cadence for Extrapolator Publication, Runtime Loops, and Verdict Consumption

3.1 Reasoning-Plane Inputs

The reasoning plane provides structured declarations:

- intent statement (task-scoped)
- tool plan and rationale
- predicted side effects
- task-conditioned allowlists (processes/domains/paths)

Session-transcript observation (reference implementation, Part II).

This detail matters because it changes the trust model of the reasoning plane. In the reference implementation, the reasoning plane is not self-declared. For OpenClaw, a separate monitor session reads other sessions' tool calls and reasoning via the platform's `sessions_list` and `sessions_history` APIs. For Cursor

Reference Runtime Stages and Adjacent Loops




Figure 4: Reference Runtime Decomposition of the Three Adjacent System-Plane Loops

and Claude Code, parallel publishers emit the same behavioral-window schema with their own `agent_type` and `agent_instance_id`. EDAMAME stores these slices, preserves contributor attribution, and merges them before correlation. This means an agent cannot retroactively edit or suppress evidence of its tool calls when the transcript or observer surface is platform-controlled. The reasoning plane therefore shifts from trust-dependent self-declaration to an independently verifiable record, though transcript integrity still depends on the platform (see Section 9, Limitations).

3.2 System-Plane Observations

The system plane consumes independent host telemetry:

- network sessions with **Layer-7 process enrichment** (binary, parent lineage, open files, command line, resource usage)
- telemetry confidence labels (`exceptions`, `anomalous`, `blacklisted`)
- host LAN visibility and local open-port observations
- posture score deltas and threat-model outputs

These telemetry labels are organized as a strict confidence hierarchy:

`sessions -> exceptions -> anomalous -> blacklisted`

The hierarchy is part of the model as a system-plane evidence-prioritization mechanism, not merely an EDAMAME implementation detail: it separates “did anything happen?” from “how strongly should we trust that this is malicious?” and allows conservative policy tuning at lower tiers. In the reference implementation, however, those labels do not replace intent correlation. The divergence / intent loop reasons primarily over raw sessions plus Layer-7, LAN, and open-port context; the vulnerability / safety-floor loop and the advisor / remediation loop consume `anomalous` and `blacklisted` labels as high-confidence operational signals.

3.3 Why Process Attribution Is Model-Critical

The two-plane model requires process attribution to remain meaningful:

- without process attribution, an allowlist collapses to destination-only checks, so benign and malicious programs contacting the same endpoint become indistinguishable
- parent-process lineage is required to detect spawn-chain manipulation, where an approved binary is launched through an untrusted wrapper chain
- open-file context is required for credential-adjacent detection, because outbound traffic alone cannot show whether sensitive files were read just before egress

This requirement is independent of implementation. Any observer that cannot attribute connections to concrete processes cannot fully realize the model.

System-Plane Telemetry Confidence Tiers




Figure 5: Signal Hierarchy with Escalating Detection Confidence




Figure 6: Layer-7 Session Enrichment Schema

Each session record is enriched with process identity, parent lineage, security signals (including `open_files` and `spawned_from_tmp`), and resource metrics. In this schema, safety-floor rules apply across the full evidence bundle, not only to a single branch.

3.4 Why MCP Matters to the Model

MCP is not only a transport convenience. It makes the model operationally portable:

1. **Structured semantics.** Typed responses reduce parsing ambiguity and let correlation logic consume machine-readable evidence directly.
2. **Transactional actions.** Execute/rollback patterns make pre-action and post-action checks composable.
3. **Framework portability.** MCP-facing tools can be consumed by OpenClaw, Cursor, Claude Code, other coding agents, and other MCP-compatible clients with minimal glue code. In the deployed contract, each producer tags its slice with agent identity and EDAMAME performs the merge centrally.

As positioned in Section 2.3, exposing continual process-attributed host telemetry as first-class MCP tools is central to this instantiation.

3.5 Multiplatform Requirement and Instantiation

Multiplatform operation is not cosmetic. Agent workflows already span Linux servers, macOS developer hosts, and Windows enterprise endpoints. This is especially true for coding agents operating across local repos, CI runners, and production deployment environments.

Accordingly, the model requires attribution beyond Linux-only eBPF paths. Our current implementation uses:

- Linux eBPF where available
- native macOS/Windows process APIs
- a socket-matching fallback when kernel telemetry is unavailable




Figure 7: Multiplatform Process Attribution Across Linux, macOS, and Windows

3.6 Observation Delay Is a Constraint, Not a Contribution

Telemetry is delayed relative to action execution. We therefore use a settling schedule (e.g., 15 s -> 30 s -> 60 s) and stop when snapshots stabilize. This is an operational accommodation, not a conceptual novelty claim. In production, that delay sits on top of the configured engine interval; demos and explicit-slice benchmarks shorten that interval, but they do not remove the underlying enrichment delay.

3.7 Two Example Rules (Simple and Realistic)

We present two example rules: one intentionally simple and one more realistic. The rules themselves are not claimed as contributions; they are concrete exam-

ples of model use and an introduction to the rule properties discussed in Section 5.

- **Example Rule A (simple, reactive):** intent-vs-observation correlation after telemetry arrives.
- **Example Rule B (more realistic, proactive+reactive):** pre-action blocking for MCP-gated operations plus Rule A for residual detection.

Example Rule A: Reactive Intent-vs-Behavior Correlation

```
process_allowed = (evidence_process in allowed_processes)
domain_allowed  = (target_ip in allowed_domains)
explained       = process_allowed AND domain_allowed

guardrail_signal = (safety_floor_findings > 0)
                  OR (vulnerability_findings > 0)

intent_mismatch  = (session_matches > 0 AND explained == false)
                  OR (credential_access AND session_matches > 0)

divergence = guardrail_signal OR intent_mismatch
```

Example Rule B: Pre-Action Guard Rule

```
action = proposed tool call (name, arguments)
intent = current declared intent

spawns_from_tmp      = action.working_dir starts_with /tmp/ | /var/tmp/ | /dev/shm/
targets_credentials  = action.arguments reference sensitive file paths
targets_blacklisted  = action.arguments contain known-blacklisted endpoints

process_in_allowlist = action.tool_binary in intent.allowed_processes
target_in_allowlist  = action.target_endpoint in intent.allowed_domains

static_violation = spawns_from_tmp OR targets_credentials OR targets_blacklisted
intent_violation = NOT (process_in_allowlist AND target_in_allowlist)

block = static_violation OR intent_violation
```

In this draft, “sensitive credential paths” is a configurable static set (for example `~/.ssh/*`, kube configs, cloud credential stores, browser-token stores, and CI token files) defined by deployment policy rather than inferred at runtime.

Rule B does not replace Rule A: actions bypassing MCP still require reactive system-plane correlation.

In this flow, the reference implementation checks blacklists, vulnerability-detection findings, and intent-conditioned allowlists within a settling-window loop. In explicit-slice benchmark mode, the behavioral model is injected

Explicit-Slice Benchmark Flow and Guardrail Reconciliation




Figure 8: Benchmark-Mode Decision Flow and Guardrail Reconciliation

immediately before this flow runs, so the figure describes the downstream correlation engine rather than the full end-to-end producer path.

Model-level multi-signal scenarios (network + identity + LAN + posture) are provided in Appendix A to keep the main narrative focused.

4. Threat Model, Assumptions, and Boundaries

We focus on runtime compromise modes that evade static checks:

- indirect prompt injection leading to unsafe tool use
- tool/MCP poisoning and capability misuse
- unexplained outbound communications during ostensibly read-only tasks
- credential-path and config-adjacent anomalies
- cross-session behavioral drift and memory poisoning

4.1 Assumptions and Why They Are Reasonable

Assumptions

- The observer is co-located with the runtime host.
- Baseline S1/S2 controls are not skipped.
- High-risk actions can require human confirmation when enabled.

Rationale

- Co-location is necessary for process attribution and parent-lineage fidelity.
- If S1/S2 are absent, runtime detection is overloaded with preventable risk.
- Human confirmation for destructive actions is standard in production ops.

Out of scope (with rationale)

- Formal properties under arbitrarily compromised reasoning planes (requires stronger mechanisms such as formal methods or trusted execution).
- Internet-scale adaptive adversarial campaigns (requires different evaluation infrastructure than the controlled benchmark used here).

4.2 Reasoning-Plane Integrity

The key question raised by reviewers is valid: if an attacker can control intent declarations, can they simply declare harmful intent and pass correlation? The answer depends on which failure case occurred:

- **Case A - Tool misuse within declared benign intent.** The model is strong. Divergence appears because observed behavior exceeds declared intent.
- **Case B - Correctly declared harmful intent.** Correlation alone is insufficient if policy allows that action. Additional governance controls (environment-aware allowlists, posture-delta thresholds, change-control invariants) are required.

- **Case C - Missing blast-radius guardrails.** If destructive production actions are broadly authorized, the root cause is authorization architecture, not telemetry correlation. The two-plane model complements S1/S2; it does not replace IAM scoping, segmentation, and approval workflows.

To raise attacker cost in Cases A/B, we apply two mitigations.

Intent-independent critical signals. The following triggers cannot be downgraded by LLM adjudication or intent declarations:

1. blacklisted destination access
2. CVE-inspired vulnerability-detection failures (gateway binding, token exfiltration, skill supply chain, sandbox exploitation)

Model-correlated signals (unexplained sessions, credential access concurrent with outbound traffic, anomalous network patterns) are promoted as divergence evidence and adjudicated by the LLM. The LLM always makes the final verdict decision, but critical signals act as non-downgradable guardrails.

Intent-declaration anomaly tracking. Motivation: compromised agents often “prepare the ground” by gradually broadening intent language before executing harmful actions. We therefore track:

1. allowlist budget expansion across redeclarations
2. abrupt intent-category shifts
3. redeclaration churn (rate spikes)

This signal is currently specified at the model level; its marginal contribution is not isolated in the Section 7 confusion-matrix metrics.

These mitigations do not remove the trust boundary entirely, but they force an attacker to simultaneously evade both declaration-level and system-level checks.

4.3 Observer Integrity and Evasion Boundaries

The model assumes a trustworthy observer process. Under kernel/root compromise, the observer itself can be tampered with. Potential hardening paths include watchdog heartbeat, cryptographic attestation, and stronger isolation.

Primary residual evasion classes are:

- **Living-off-the-land** (allowed process + allowed destination)
- **timing attacks** (harm before first informative snapshot)
- **observer evasion/tampering**
- **credential-less exfiltration** (transient file access not captured in time)

The most operationally important combined risk is living-off-the-land plus timing. This combination is dangerous because behavior can remain syntactically “allowed” while executing before the most informative telemetry fields are available. Near-term mitigations are shorter enrichment intervals, sticky process-context retention, stricter defaults for high-risk tasks, and proactive Rule B

gating for MCP-mediated operations. We detail mitigation priorities and benchmark blind spots in Section 9.

4.4 State Isolation and Observer Independence

A subtle but critical attack surface in agentic frameworks is *shared mutable state*. Many frameworks persist inter-session context – behavioral models, verdicts, operational memory – in flat files within the agent’s workspace (e.g., `MEMORY.md`, `SOUL.md`). These files are readable and writable by any process on the host. This creates multiple attack vectors:

1. **Model poisoning.** A compromised process writes false behavioral predictions into the shared memory file, causing the detection layer to whitelist malicious traffic patterns as “expected.”
2. **Verdict tampering.** A compromised process overwrites the stored verdict (e.g., writing `CLEAN` over a `DIVERGENCE` determination), suppressing alerts before they reach human operators.
3. **Self-covering reasoning plane.** A compromised agent can edit its own memory to erase evidence of divergent behavior, defeating any retrospective audit.

These attacks are particularly dangerous because they require no privilege escalation – only filesystem write access to a user-owned directory.

Our architecture mitigates this by *internalizing detection state*. The divergence engine runs inside the EDAMAME host runtime: the EDAMAME Posture daemon in the reference benchmark deployment, and the EDAMAME Security app on workstation-hosted Cursor paths. Behavioral models, verdict history, and engine configuration live in process memory and are accessible only through authenticated MCP endpoints (pre-shared key). Multiple reasoning-plane producers can push agent-tagged slices via `upsert_behavioral_model`; EDAMAME merges those slices into one observer-owned correlation window while retaining contributor attribution metadata. The reasoning-plane agent can contribute predictions, but cannot read or modify verdicts, engine state, or critical-signal rules.

This separation provides three concrete assurances:

- **Unilateral tamper resistance.** Compromising the reasoning plane alone is insufficient to suppress detection. The observer’s state is unreachable from the agent’s process or filesystem.
- **Independent alerting.** The observer-owned notification pipeline (Slack, Telegram, webhook) is shared across agentic and divergence/vulnerability outputs, not owned by the agent. Even a fully compromised reasoning plane cannot intercept or cancel outbound alerts.
- **Critical-signal invariance.** Intent-independent rules (Section 4.2) are evaluated inside the observer. No external input can downgrade a critical-signal finding.

The extrapolator skill retains OpenClaw’s memory system for its own operational checkpoint (a cursor tracking the last-processed session). This is acceptable: the checkpoint is not security state. Corrupting it causes the extrapolator to reprocess sessions but cannot influence EDAMAME’s verdicts.

5. Properties of the Decision Rules

This section can be read as a short extension of Section 4 (Threat Model): it states properties of the deterministic decision rules (Example Rules A-B) that hold by construction. These are not deep theorems; they are simple, auditable invariants meant to make system behavior transparent. We explicitly distinguish these from stronger formal properties (PAC bounds, temporal SMT constraints, cryptographic proofs) provided by systems such as Pro2Guard, Agent-C, and VET (Section 2.2). Our contribution is practical, not formal.

5.1 Definitions

For each step t :

- I_t : declared intent (reasoning plane)
- C_t : context constraints derived from I_t (e.g., allowlists)
- O_t : system-plane observation snapshot after settling
- $K(O_t)$: critical-signal predicate (true if blacklisted traffic or vulnerability-detection failure)
- $E(I_t, C_t, O_t)$: “explained” predicate (true if O_t is consistent with intent)

The two-plane alert predicate is:

$$D(I_t, C_t, O_t) = K(O_t) \vee (O_t \neq \emptyset \wedge \neg E(I_t, C_t, O_t))$$

5.2 Property 1 (False-Positive Monotonicity)

For any anomaly predicate A over observations, let $D_gated = A$ AND (NOT E) and $D_critical = K$ OR D_gated . Then for benign runs where K does not fire:

$$P(D_critical = 1 \mid \text{benign}, K=0) = P(D_gated = 1 \mid \text{benign}) \leq P(A = 1 \mid \text{benign})$$

This follows from set inclusion: $\{D_gated=1\}$ implies $\{A=1\}$. Intent-gating cannot increase false positives relative to anomaly-only alerting when no critical signal is triggered.

Remark (critical-signal false positives). Critical signals *can* produce false positives for benign operations that coincidentally match (e.g., a legitimate connection to a recently-blacklisted destination).

5.3 Property 2 (Settling Window and Miss Rate)

Let L be the time until an action’s effects become observable. With maximum wait W : $P(\text{miss}) = P(L > W)$. Progressive re-checks (sampling at $t_1 < t_2 < \dots < W$) do not change this bound but reduce the expected detection latency when $L < W$.

5.4 Property 3 (Example Rule A Correctness Under Complete Telemetry)

Assume the evidence bundle is complete: `session_matches > 0` iff the endpoint was contacted by `evidence_process` after the watermark. Example Rule A returns `divergence=true` iff the observed egress is not authorized by intent OR a critical signal fires. This follows directly from the algorithm’s structure and is not a claim about real-world completeness. In practice, this assumption is violated by settling windows and refresh delays (Section 3.6, Section 9.1), so Property 3 should be read as a rule-level sanity property, not an empirical coverage claim.

5.5 Property 4 (Pre-Action Blocking is Conservative)

Example Rule B blocks on `static_violation OR intent_violation`. Since both predicates are computed from the proposed action and current intent (no telemetry), blocking is immediate (zero latency). The cost is over-blocking: legitimate actions that happen to match a floor pattern (e.g., a build tool that writes to `/tmp/`) require human confirmation.

5.6 Scope of These Properties

These properties characterize the decision rules under idealized assumptions (complete telemetry, correct allowlists, honest intent for Property 1). They do not provide:

- Safety properties under partial telemetry (missed sessions, eBPF gaps)
- Bounds on evasion probability (see Section 4.3)
- Temporal safety properties (no ordering constraints on actions)
- Probabilistic properties on detection rates (the Wilson CI in Section 7.4 is empirical, not analytical)

Systems that require these stronger properties should consider Pro2Guard (PAC bounds), Agent-C (temporal SMT), or VET (cryptographic attestation).

Part I takeaway: the model provides an auditable way to correlate declared intent with host-observed behavior, but by itself it does not provide zero-latency interception, complete living-off-the-land coverage, or formal safety proofs. Part II therefore focuses on implementation feasibility and measured operational behavior.

Part II — Implementation and Evaluation

Part II demonstrates feasibility and relevance: Section 6 describes our implementation, Section 7 reports evaluation results, Section 8 documents reproducibility assets, Section 9 states limitations and near-term work, and Section 10 discusses ethics and safe use.

6. Implementation (OpenClaw + EDAMAME Posture)

Our reference implementation combines one OpenClaw reasoning skill with EDAMAME-internal runtime loops. We distinguish the architecture itself from the benchmark mode used later in Section 7.

Reasoning-plane publication:

- **skill/edamame-extrapolator/** – Reads raw session history via `sessions_list` and `sessions_history`, distills a compact behavioral model, and publishes it via the behavioral-model upsert API. Production cadence is `*/5`; tests and demos accelerate this to `*/2`.
- Workstation integrations such as Cursor can forward raw reasoning sessions through the raw-session publication API, allowing EDAMAME to generate the contributor slice internally before merge.

System-plane evaluation inside EDAMAME:

- **Internal divergence engine** (inside EDAMAME Posture) – Performs the downstream correlation loop. The production default interval is 300 s, but the CLI can shorten it; executive demos use 60 s and the explicit-slice benchmark helper can drive it faster because it bypasses producer cadence. Each cycle proceeds in three stages:
 1. **Deterministic evidence stage** (Rust, `divergence_engine.rs`): Stateless pure functions correlate behavioral model predictions against live telemetry (sessions, anomalous, blacklisted, exceptions) and typed contextual state. The same local stage also accumulates intent-relative evidence for the final verdict path.
 2. **Optional LLM adjudication stage**: The deterministic evidence bundle can be submitted to the configured provider for soft-signal classification and operator-facing alert synthesis. The provider returns a structured `LlmDivergenceDecision` plus an optional `LlmHumanAlert`.
 3. **Guardrail reconciliation** (`reconcile_with_guardrails`): Critical signals – especially vulnerability-detection findings and blacklisted-traffic evidence – are non-downgradable. If the LLM attempts to override a CRITICAL deterministic finding, the guardrail forces DIVERGENCE with provenance `GUARDRAIL_FORCED_DETERMINISTIC`. Otherwise the soft-signal verdict is accepted with provenance such as `LLM_CONFIRMED` or `LLM_OVERRULED_SOFT_SIGNALS`. The local deterministic stage al-

ways runs first. Optional LLM output can influence soft-signal classification and operator wording, but it cannot downgrade CRITICAL deterministic findings. If the behavioral model is stale (>20 min), the engine produces a STALE verdict. Verdicts are accessible via `get_divergence_verdict`.

- **Internal vulnerability / safety-floor loop** – Evaluates model-independent runtime incident checks (for example gateway binding, token exfiltration, skill supply chain, sandbox exploitation, `/tmp` lineage, and credential-file plus egress combinations) against the same system-plane substrate.
- **EDAMAME posture/remediation loop** (internal) – Reviews advisor todos and executes safe fixes through the native agentic pipeline. In OpenClaw, `skill/edamame-posture/` acts as a thin MCP facade for invoking these capabilities; scheduled remediation ownership remains inside EDAMAME.

Benchmark interpretation for this draft:

- The aggregate metrics in Section 7 come from explicit-slice benchmark mode. The harness derives a scenario-specific behavioral model and injects it with `upsert_behavioral_model`, so the quantitative loop measures downstream correlation behavior under accelerated cadence rather than transcript-to-model extraction accuracy.

Supporting components:

- OpenClaw plugin bridging to MCP: `extensions/edamame-mcp/`
- Live benchmark harness: `tests/benchmark/run_live_suite.py`
- Metric summarizer: `tests/benchmark/summarize_results.sh`
- Canonical claim binding: `docs/CLAIM_ARTIFACT_INDEX.md`
- Public paper source and generated publication bundle: `agent_security`
- OpenClaw workstation/package repository used for the OpenClaw path: `edamame_openclaw`
- Cursor workstation/package repository used for the Cursor path: `edamame_cursor`

Behavioral-model and verdict state are internal to the EDAMAME divergence engine. OpenClaw memory remains optional operational checkpointing for the extrapolator cursor and is not authoritative detection state.

6.1 MCP Integration Path

EDAMAME Posture exposes a local, PSK-authenticated MCP endpoint via Streamable HTTP (`http://127.0.0.1:3000/mcp`). OpenClaw calls EDAMAME MCP tools through a local OpenClaw plugin (`extensions/edamame-mcp`) that implements a minimal HTTP MCP client and exposes EDAMAME tools as native OpenClaw tools:

OpenClaw (native tool call) ->

edamame-mcp plugin ->
EDAMAME MCP (HTTP) ->
tool result

MCP Integration Notes Section 3.4 explains why MCP is model-relevant. At implementation level, the key point is that OpenClaw can consume EDAMAME as native tools (typed JSON over local PSK-authenticated HTTP), so correlation logic can run without shell parsing or privileged subprocess glue.

6.2 Trace-Backed Evidence Bundles

The live benchmark suite records:

- baseline and post-action telemetry snapshots
- inject logs with VM-side timestamps
- MCP call outputs (todos, action history)
- a run manifest with scenario set hash and git SHA

These artifacts enable independent inspection of every verdict.

Two engineering lessons shaped the current draft. First, separating producer cadence from verdict cadence made the downstream correlation loop much easier to explain and debug. Second, benign background traffic proved harder to stabilize than the core lineage or credential-theft scenarios. That is why this revision treats the current snapshot as useful but provisional rather than as a polished production operating point.

6.3 Claim Gating and Validation Maturity Tiers

To avoid confusion between runtime **evidence bundles** (system telemetry used by the divergence engine) and paper-level **validation status** (what this manuscript has empirically established), we use the following maturity tiers:

- T0: implemented
- T1: demonstrated in integration tests
- T2: measured in trace-backed benchmark artifacts
- T3: independently replicated

Automation blocks unsupported strong-claim language unless minimum tier floors are met. Published metric claims are anchored to a single manifest/summary lineage in `docs/CLAIM_ARTIFACT_INDEX.md`, so run IDs, hashes, and benchmark modes remain auditable.

6.4 Signal Coverage vs Validation Tier (What Is Empirically Supported)

The system targets multi-signal attacks, but empirical support is non-uniform. Table 1 summarizes current status:

Channel	Tier	Notes
Network egress (sessions, anomalous, blacklisted)	T2	Trace-backed live suite (paper metrics)
Process lineage (parent_pid, parent_name, parent_path)	T1	Layer-7 enrichment integration-tested via MCP
Open file access (per-process file descriptors)	T1	Credential-theft detection via open_files
Exceptions / policy violations	T1	Sanity-checked, no dedicated T2 protocol
MCP wiring (advisor_get_todos, action_history)	T1	OpenClaw -> plugin -> MCP path verified
Posture scoring (get_score, advisor_get_todos)	T1	Triage/demo flows; not yet T2-measured
Vulnerability detection (config/env safety scans)	T1	Unsafe configuration plus host-open-port gateway exposure checks
Cross-session drift (memory poisoning)	T0/T1	Partially exercised; not in T2 protocol

6.5 Operational Context: Posture Scoring and Compliance

Beyond individual network signals, EDAMAME Posture continually computes a composite security score (0-100) with subcategory breakdowns (network, system integrity, system services, applications, credentials) and maps the endpoint state against compliance frameworks (CIS Benchmark Level 1/2, ISO 27001/2, SOC 2). This provides two advantages for the two-plane model:

1. **Baseline context.** Endpoint hardening level (e.g., 90/100 vs 45/100) changes risk posture and response aggressiveness.
2. **Score-delta detection.** A posture drop during a task (e.g., 65 to 50) is an independent system-plane signal, even when individual sessions look benign.

These capabilities are exercised in the extrapolator/internal divergence engine flow and demo path (T1). Their incremental impact on T2 live detection accuracy remains future work. Even at current maturity, this context already improves operator triage and supports safer action review in T1-validated workflows.

7. Evaluation: The BadAgentUse Benchmark

We use **BadAgentUse** as a live, trace-backed benchmark for the downstream correlation loop against real endpoint telemetry. Unlike prior benchmarks that simulate network traces or rely only on sandboxed logs, BadAgentUse injects live traffic into an instrumented Linux VM and scores detection against actual EDAMAME Posture session captures. The current aggregate numbers should not be read as full end-to-end OpenClaw/Cursor deployment scores.

7.1 Setup

- **Environment:** Lima VM, Ubuntu 24.04, 4 vCPU / 4 GiB
- **Observer:** EDAMAME Posture daemon (eBPF-backed on aarch64 Linux) with MCP Streamable HTTP endpoint on port 3000
- **Reasoning-plane path:** In production, OpenClaw publishes behavioral-model slices through the `edamame-extrapolator` skill, and Cursor publishes through the same EDAMAME contributor contract. For the quantitative BadAgentUse runs reported here, however, the harness injects scenario-specific behavioral-model slices directly via `upsert_behavioral_model` after deriving them from the scenario contract and live baseline telemetry. This isolates the downstream divergence engine from producer-side cron cadence and transcript-to-model non-determinism. End-to-end extrapolator behavior is validated separately by integration and E2E tests rather than by the main aggregate metric loop.
- **Cadence for the measurements in this draft:** production defaults use 300 s internal divergence/vulnerability intervals and $\ast/5$ OpenClaw publication/reader cron. Executive demos use 60 s internal loops and $\ast/2$ OpenClaw cron. The explicit-slice benchmark helper accelerates the divergence engine further (down to 1 s in the helper path) because the behavioral slice is injected directly before verdict polling.
- **Protocol:** 50 versioned JSON scenario files in the benchmark scenario directory (`tests/benchmark/live-scenarios`)
- **Iterations:** configurable via `--iterations`; the current reproducible branch snapshot is a 10-iteration live sweep over the current 50-scenario set
- **Seeds:** 10 distinct seeds in the current branch snapshot, with deterministic target-IP rotation
- **Trial unit:** one scenario evaluated per seed-iteration pair. The confusion matrix (Section 7.4) reports $n=468$ valid runs across 10 seeds, with 32 skipped runs preserved in the readiness scorecard. Each iteration varies seeds and settling behavior; all iteration traces remain in artifacts.

Protocol Parameters

Parameter	Value	Notes
Settling schedule	15 s $\rightarrow$ 30 s $\rightarrow$ 60 s	Exponential back-off; max wait W = 60 s
Stability criterion	2 consecutive stable snapshots	Snapshot = sessions + Layer-7 + posture unchanged
Layer-7 full refresh	300 s (5 min)	Process metadata, parent lineage, open files
Sensitive-file scan	30 s	Credential paths only; sticky retention
Anomaly-score threshold	0 (any non-zero triggers)	Conservative; tunable per deployment
Posture-delta threshold	≥ 1 point drop	Monitoring signal (not hard-gate)
Blacklist sources	FireHOL Level 1 + custom	Pinned at VM provision time
Allowlist semantics	Exact domain OR IP; no wildcards	AND logic: process AND destination must match
MCP endpoint	127.0.0.1:3000, PSK-authenticated	Localhost-only by default
LLM (attack detection)	openai/gpt-5.1	Deterministic tiers do not use LLM; compatible Azure-hosted OpenAI Responses models are supported as alternate providers

The important timing distinction is that the 51.6 s median reported later combines accelerated engine cadence with slower underlying telemetry enrichment; it is not evidence that the production default timer itself runs on a 51.6 s interval.

7.2 Scenario Design (Versioned Categories)

In BadAgentUse, a **scenario** is a versioned JSON test contract that specifies: (1) declared intent context, (2) injected runtime action or traffic pattern, (3) expected enforcement label (DIVERGENCE, CLEAN, or DESIGNED_MISS), and (4) the telemetry checks needed for scoring.

BadAgentUse covers eight categories spanning five signal families:

Cat.	N	Signal	Description
A. Network-Intent	17	Network	Undeclared egress, partial-allowlist AND, lateral movement
B. Process Lineage	6	Layer-7	/tmp/ scripts, interpreter wrap, nested chains
C. Credential Theft	7	Layer-7 files	SSH, kube, gnupg, browser cookies, combined exfil
D. CVE Vulnerability Detection	6	Layer-7+allow	gateway_binding / token_exfil / credential overrides
E. Posture & Policy	–	Posture	Score, policy, threats, HIBP (integration checks) [†]
F. LAN Awareness	–	LAN	Device discovery, scan, anomaly APIs (integration checks) [†]
G. Adversarial	8	Mixed	LotL, timing, burst, path substitution
H. Noisy Benign	6	Network	Clean traffic, DNS, NTP, auto-update

[†] Categories E and F are validated via dedicated integration scripts (Section 7.7), not as JSON benchmark scenarios. The 50 benchmark scenarios span categories A–D, G, H.

Enforcement labels. Each scenario carries an `expected_enforcement` field: DIVERGENCE (39 scenarios), CLEAN (8 scenarios), or DESIGNED_MISS (3 scenarios). The DESIGNED_MISS label explicitly marks living-off-the-land attacks that Example Rule A cannot detect by construction (Section 4.3).

7.3 Metrics

We adopt the metric framework from AgentSentinel [Hu et al., 2025]:

- **DSR** (Detection Success Rate) = $TP / (TP + FN)$ – equivalent to recall for attack scenarios
- **FPR** (False Positive Rate) = $FP / (FP + TN)$ – rate of benign traffic incorrectly flagged
- **FNR** (False Negative Rate) = $FN / (FN + TP)$ – attacks that evaded detection
- **ASR** (Attacker Success Rate) = FNR – attacks that succeeded despite monitoring
- **FT** (False Trigger) = FPR – identical to FPR for symmetry
- **TTD** (Time-to-Detect) = median and p95 detection latency (ms), measured by the harness from injection timestamp to the first scored detection event

- **Wilson CI95** on precision and recall for the current snapshot. Because the benign denominator is still small and being rebuilt, we discuss FPR qualitatively rather than treat its current interval as stable evidence.

All metrics are computed from trace-backed NDJSON results by the benchmark summarizer, which also produces per-category and per-seed stability breakdowns.

7.4 Statistical Confidence Reporting (Wilson Interval)

To bound measured proportions in the current branch snapshot, we use the Wilson score interval (Wilson, 1927). For k successes out of n trials and $z=1.96$ (95% confidence):

```
p_hat = k/n
den    = 1 + z^2/n
center= (p_hat + z^2/(2n)) / den
half   = z * sqrt(p_hat*(1-p_hat)/n + z^2/(4n^2)) / den
CI95   = [center - half, center + half]
```

In the current reproducible branch snapshot used by this draft:

	Predicted DIVERGENCE	Predicted CLEAN
Ground-truth attack	TP = 317	FN = 31
Ground-truth benign	FP = 0	TN = 120

Precision = $317/317 = 1.000$ and recall = $317/348 = 0.911$, yielding CI95 [0.988, 1.000] and [0.876, 0.937], respectively ($n = 468$ valid runs, 10 seeds). With $TN=120$ and $FP=0$, the measured benign slice in this sweep stayed clean, but we still treat benign generalization as environment-dependent outside the versioned scenario contracts used here.

The 31 FN are not uniformly distributed. Thirty come from the three scenario contracts whose `expected_enforcement` is `DESIGNED_MISS`; these scenarios are intentionally retained as attack-class rows in the confusion matrix so that the reported recall continues to reflect known blind spots rather than hiding them. The remaining miss is a single non-designed miss in network-intent-correlation. The aggregate result therefore reflects one residual accidental miss plus an explicitly versioned allowlist-respecting living-off-the-land blind-spot bucket, not broad degradation across the suite.

7.5 Results

Current Reproducible Branch Snapshot *Results below come from the current reproducible branch snapshot in `artifacts/live-paper-summary.json` (50 planned scenarios; 468 valid runs across 10 seeds, with 32 skipped runs retained in the readiness scorecard). These aggregate metrics measure*

EDAMAME's downstream correlation loop under the explicit-slice producer mode (`upsert_behavioral_model`) on a live runtime, not the full end-to-end transcript-publication path. To keep the sweep tractable, the benchmark uses accelerated cadence plus benchmark-scoped runtime assists recorded in the manifest, including score-kick flushing and preservation of benchmark-provided `ingested_at` timestamps for slices tagged `benchmark`; these choices reduce waiting and verdict-latency overhead but do not change scenario labels, confusion-matrix accounting, or post-hoc scoring. Cursor and Claude Code use the raw-session producer mode used by workstation integrations and are validated separately rather than as a second quantitative cohort in this draft. The snapshot below is the completed live sweep for the current scenario set, not a partial provisional subset.

Snapshot metadata:

- `run_id`: run-live-20260313T194140Z-live
- `git_sha`: aca9156
- `scenario_set_version`: fa1266a7eb7b2645... (full hash in docs/CLAIM_ARTIFACT_INDEX.md)
- `mode`: live
- `benchmark_mode`: live

Metric	Value
Total runs	468
Precision	1
Recall	0.9109195402298851
Precision CI95	[0.988026490330431, 1]
Recall CI95	[0.8763311151407897, 0.9365346744645328]
Median latency (ms)	51590
p95 latency (ms)	66879
Rollback reliability	1
Stability (precision stddev)	0
Stability (recall stddev)	0.009001831155373185

Readiness note. The current committed lineage passes the repo's publication gate: `artifacts/readiness-scorecard.json` is G0 because the precision CI95 lower bound is 0.988, the recall CI95 lower bound is 0.876, and the skipped-run ratio is 6.4%.

Per-Category Detection




Figure 9: Canonical Live Benchmark Summary for the Current Live Sweep

Category	Run N	DSR	FPR	Notes
Network intent correlation	160	0.990	0.000	one remaining miss in a partial-allowlist Python case; benign allowance cases stayed clean
Process lineage verification	59	1.000	—	perfect across tmp-spawn, wrapper lineage, and nested chains
Credential theft detection	30	1.000	—	SSH, kube, gnupg, browser-cookie access patterns
Credential-network combined	10	1.000	—	combined credential + egress behavior

Category	Run N	DSR	FPR	Notes
Intent-proof floor	59	1.000	–	floor-policy and temporary-spawn floor checks all hit in this sweep
Lateral movement	10	1.000	–	nc-based lateral-probe preparation behavior
Adversarial evasion	80	0.625	–	remaining misses are concentrated in allowlist-respecting LotL/timing stress cases
Noisy benign	60	–	0.000	benign DNS, NTP, and auto-update traffic stayed clean in this sweep

Categories E (Posture & Policy) and F (LAN Awareness) are validated via dedicated integration scripts and summarized separately in Sections 7.7 and 7.8; they are not folded into the confusion matrix above.

Analysis Observed false negatives (31). The miss surface is concentrated and largely intentional. Thirty of the 31 attack-side misses come from the three `DESIGNED_MISS` scenario contracts repeated across 10 seeds; those rows remain in the attack class and therefore still lower the reported recall. The remaining miss is one non-designed miss in network-intent-correlation. The practical implication is that this branch snapshot remains strong on process-lineage, credential-access, credential-plus-egress, intent-proof-floor, lateral-movement, and most core network-intent cases, while the dominant unresolved weakness is still allowlist-respecting living-off-the-land behavior rather than a broad loss of detector coverage.

Benign stability in the measured set. No false positives were observed in this 10-seed sweep: noisy-benign scenarios stayed clean and benign allowlist cases no longer produced the earlier operator-fatigue pattern. That is an important publication result for this branch snapshot, even though real deployments still need environment-specific allowlisting and baseline review.

High-level takeaway: the current 10-seed live evidence set reinforces the system’s value as a runtime correlation layer for lineage, credential access, lateral movement, and undeclared egress, while sharpening one main caveat: adversarial/LotL behavior that remains inside expected tools or destinations is still the dominant recall bottleneck. The system is therefore best understood as a practical early-warning layer, not as an interception-grade prevention mechanism.

7.6 Ablation Studies

To interpret the overall Section 7.5 result, we run ablations that isolate which components produce most of the detection lift and where misses originate. The benchmark runner supports three ablation modes via `--ablation`:

Ablation	Disabled	Expected Impact
no-introspection (ablation flag name)	CVE vulnerability detection	Cat. D degrades to CLEAN
no-l7	Layer-7 enrichment	Cat. B, C lose detection
no-settling	Settling window	Transient misses; timing evasion

These ablations define component contributions. Running each mode writes comparative summaries to `artifacts/live-paper-summary.json` for the benchmark-summary visualization used in companion materials. In this draft, full `no-floor` / `no-l7` / `no-settling` comparative tables are not yet included in the current confusion-matrix snapshot. We therefore treat this subsection as implementation status rather than as a measured comparative claim set.

eBPF vs netstat2 fallback. `tests/test_l7_ebpf_ablation.sh` runs the same Layer-7 checks twice: with eBPF kprobes (default on kernels ≥ 5.8), then with tracefs unmounted to force netstat2 `/proc/net/*` fallback.

Mode	Checks	Pass	Layer-7 Latency	Sessions
eBPF (kprobe)	9	9/9	15 s	50
netstat2 (fallback)	9	9/9	45 s	31

Both modes pass 9/9 and preserve the same core enrichment fields: PID, process name, parent lineage, temp-spawn signal, open-file context, cwd, memory, and CPU usage. Latency differs (15 s with eBPF vs 45 s fallback) because eBPF only accelerates socket-to-PID mapping. After PID resolution, both paths share the same `/proc` enrichment logic. Fallback can miss very short-lived sockets between polls, but for connections that survive at least one poll cycle, detection capability is preserved.

7.7 Multi-Signal Observer Validation

After isolating component-level effects, we verify observer breadth: the system plane is not only a network monitor, but a multi-signal telemetry surface. Categories E (Posture) and F (LAN) validate that the system-plane observer extends beyond network sessions. The test suite verifies MCP tool accessibility for five signal families:

1. **Network sessions:** session and exception APIs (Categories A-D, G-H)
2. **Anomaly/blacklist:** anomalous-session and blacklist APIs
3. **Posture:** score, policy, and advisor APIs (Category E)
4. **LAN:** device, host-port, and auto-scan APIs (Category F)
5. **Identity:** pwned-email and breach APIs (Category E)

For benchmark/tool-access validation, the suite calls the session API; the internal divergence and vulnerability loops themselves consume current-session snapshots within the EDAMAME host runtime.

All five families are exercised through deterministic tool-access checks (no LLM in this subsection), confirming that the telemetry surface spans these dimensions without infrastructure changes. The broader pass/fail state of those checks is summarized in Section 7.8 rather than collapsed into a single 100% claim.

7.8 Empirical Validation: Four-Tier Test Architecture

Beyond the benchmark suite, we maintain a four-tier test architecture:

Tier 0: Smoke and component validation. Smoke, component, extrapolator-validation, and posture scripts cover daemon health, MCP tool visibility, session APIs, behavioral model engine handoff, and posture skill runtime. 79+ checks across 4 test groups, all passing.

Tier 1: Layer-7 telemetry validation. L7 telemetry and eBPF/fallback ablation scripts verify parent lineage, open-files detection, field correctness, and coverage (20 checks plus 2x9 eBPF/fallback checks).

Tier 2: Attack and policy/LAN detection. Attack-detection, posture-policy, and LAN-awareness scripts exercise the hybrid Rust-deterministic + LLM-adjudication pipeline for attack detection, alongside deterministic policy and LAN checks to validate posture/todos and LAN/anomaly/blacklist/HIBP APIs.

Tier 3: Two-plane E2E and vulnerability detection. Two-plane E2E and vulnerability-detection scripts validate the full worker + extrapolator + inject + detector pipeline plus CVE-inspired checks across token exfiltration, exposed gateways, weaponized skills, sandbox escapes, credential sprawl, and memory poisoning.

Tier	Script	Checks	LLM?	Pass Rate	Notes
0	PoC suite	29	Mixed	97% (28/28 + 1 skip)	Infrastructure, telemetry, MCP tools, agent
0	Smoke	14	No	100%	Daemon health, eBPF kprobes
0	Skill components	55	No	100%	MCP tools, engine handoff
0	Extrapolator validation	–	Yes	100%	Behavioral model pipeline
0	Posture runtime	24	Yes	100%	Skill runtime invocation
1	L7 telemetry	20	No	100%*	PID, process name, cmd attribution
1	L7 eBPF ablation	2x9	No	Partial	eBPF vs fallback comparison
2	Attack detection	22	Mixed	Partial**	Engine verdict pipeline under stabilization
2	Posture policy	8	No	87%	7/8 pass
2	LAN awareness	6	No	83%	5/6 pass
3	Two-plane E2E	–	Yes	Partial**	Worker + extrapolator + detector pipeline
3	Vulnerability detection	66	Yes	86%	57/66 pass

*L7 Tier 1: Basic L7 attribution (PID, process_name, full command line) passes consistently with eBPF kprobe programs active (`track_connect_v4/v6`, `trace_udp_send`, `trace_sendto`). Extended enrichment fields (parent_process_path, open_files) are populated asynchronously by the L7 enricher and depend on /proc scan depth; these fields are not tested in the core validation suite since they are additive context for the divergence engine rather than structural correctness requirements.

**Engine verdict: The hybrid Rust/LLM divergence engine returns structured verdicts only after a behavioral model has been ingested and the 60-second settling window has elapsed. The `test_poc.sh` suite correctly handles this as a SKIP condition when no agentic tick has occurred. Tier 0 component tests (55/55) confirm all engine MCP tools are reachable.

7.9 Discussion

Latency is telemetry settling, not model delay. The observed 40-60s TTD is dominated by EDAMAME session aggregation and Layer-7 refresh cadence. Without settling-window checks (Section 3.6), many runs would be scored on stale data.

Process attribution is load-bearing. Attack and benign runs can generate identical traffic to the same endpoint; only process attribution separates `curl` from `python3` at the same IP. Without Layer-7 enrichment, separation collapses.

Claim discipline is mandatory. We do not claim comprehensive coverage. BadAgentUse validates five signal families and explicitly labels three designed blind spots. The suite is extensible via new JSON scenarios, without code changes.

Benign calibration is improved in this snapshot, not solved forever. The current measured sweep reports FPR=0.000 for the noisy-benign category,

so the earlier operator-fatigue pattern is not present in this branch snapshot. That does not eliminate future benign-noise burden in deployment: local update traffic, private mirrors, and organization-specific tooling can still require environment-specific allowlisting and baseline review.

Comparative baselines are currently absent. BadAgentUse does not yet include head-to-head comparisons against AgentSentinel, AgentArmor, LlamaFirewall, Pro2Guard, or Agent-C. This benchmark validates internal correctness of the correlation protocol and telemetry surface, not superiority. Direct comparison requires adapting either the harness for interception-style monitors or those systems for MCP telemetry (Section 9.2).

7.10 Observer Resource Footprint

We measured EDAMAME Posture’s CPU and memory footprint on the same 4 vCPU / 4 GiB Lima VM used for the benchmark evaluation (Ubuntu 24.04, kernel 6.8.0-106, aarch64, edamame-posture 1.1.4 installed via APT). Sampling used `pidstat` at 5 s intervals over 60 s per phase, with `/proc/<pid>/status` for memory snapshots. Three phases were measured: **idle** (daemon running with eBPF-backed packet capture but no behavioral model or divergence engine activity), **active** (behavioral model pushed, divergence engine correlating), and **load** (active engine plus concurrent benign HTTPS traffic to multiple external hosts).

Phase	CPU % (VM)	CPU %		VmSize (MiB)	Threads
		(bare-metal est.)	RSS (MiB)		
Idle	0.83	0.79	497	1932	16
Active	0.55	0.52	498	1932	16
Load	1.37	1.30	499	1932	16

The VM runs under Lima 2.0.3 with Apple Virtualization.framework (VZ) on macOS/aarch64. VZ uses a thin hypervisor with hardware-assisted paravirtual I/O and direct memory mapping (no balloon driver). Published benchmarks and Apple developer documentation place VZ CPU overhead at 3–8% for compute-bound workloads; memory overhead is negligible because guest pages map directly to host physical memory without a software translation layer. We apply a conservative 5% correction (reported VM CPU% $\times$ 0.95) for the bare-metal estimates in the table. This is an extrapolation, not a measured bare-metal result; readers should treat the VM column as the primary measurement.

For context, widely deployed Linux EDR agents report comparable or higher baseline footprints: CrowdStrike Falcon sensor typically uses 1–5% CPU and 200–500 MiB RSS depending on policy and telemetry volume; Elastic Agent reports 150–300 MiB RSS at idle. EDAMAME Posture’s ~500 MiB RSS includes

the eBPF-backed network session capture engine (flodbadd), Layer-7 process attribution, the agentic divergence engine with LLM integration, posture scoring, and the MCP server – a broader functional surface than a pure detection agent. Virtual memory (VmSize) is dominated by the Rust async runtime pre-allocated arena and is not indicative of physical memory pressure.

Artifact: `artifacts/footprint-summary.json`; reproduction script: `scripts/measure_footprint.sh`.

8. Reproducibility

All source code, benchmark scenarios, harness scripts, and canonical run artifacts are contained in the public `agent_security` repository. Key artifact paths:

- `tests/benchmark/live-scenarios/`: 50 versioned JSON scenario contracts
 - `tests/benchmark/run_live_suite.py`: benchmark harness (orchestrates live runs inside a Lima VM)
 - `tests/benchmark/summarize_results.sh`: metric computation from NDJSON traces
 - `tests/benchmark/validate_claims.sh`: checks that paper claims stay within evidence bounds
 - `docs/manifest-schema.json`: formal JSON Schema for the run manifest
 - `docs/CLAIM_ARTIFACT_INDEX.md`: canonical claim-to-artifact lineage binding
- Attack-injection scripts for executive demos live in the `edamame_openclaw` repository (see also the private `openclaw_security` monorepo).

Canonical run artifacts shipped with this draft:

- `artifacts/live-paper-results.ndjson` – per-scenario NDJSON traces
- `artifacts/live-paper-summary.json` – aggregate metrics and confusion matrix
- `artifacts/live-paper-manifest.json` – full run manifest (scenario groups, versions, cron config)
- `artifacts/paper-bundle/*` – generated whitepaper bundle with embedded metrics

Single-command artifact-level reproduction (recomputes the summary from the committed NDJSON without requiring a live VM):

```
bash scripts/reproduce_results.sh
```

Full live reproduction (requires a configured Lima VM with EDAMAME Posture):

```
python3 tests/benchmark/run_live_suite.py --mode live
```

9. Limitations and Future Work

9.1 Current Limitations

1. **Living-off-the-land blind spot.** The primary residual risk is LotL behavior (Section 4.3): allowed processes exfiltrating to allowed destinations. BadAgentUse explicitly labels 3 of 50 scenarios as `DESIGNED_MISS` for this reason. Closing this gap requires behavioral baselines, anomaly fusion, or payload-aware inspection.
2. **Layer-7 refresh latency.** Full Layer-7 enrichment uses a 5-minute production default (sensitive-file scan every 30 s on Linux). Demo and benchmark runs shorten the engine interval, but they do not eliminate this underlying enrichment delay. Correlation rules that rely on `spawned_from_tmp` or `open_files` can miss short-lived processes if snapshots are stale. The `no-settling` ablation (Section 7.6) quantifies this.
3. **Reactive detection delay (operationally significant).** Example Rule A is post-observation and therefore delayed. The 51.6 s median reported in Section 7 comes from an accelerated benchmark configuration, not the steady-state production default. In high-consequence environments, even that shortened window may still allow harmful actions before alerting. Example Rule B reduces this for MCP-gated actions, but direct shell/library paths can bypass pre-action checks.
4. **Interception trade-off vs AgentSentinel.** AgentSentinel reduces timing exposure for intercepted operations by pausing execution inline; our approach prioritizes portability and lower integration friction. This is a real design trade-off, not a superiority claim.
5. **Governance boundary (Cases B/C).** If the reasoning plane truthfully declares a harmful action and policy allows it, correlation alone is insufficient (Section 4.2). Authorization architecture, change control, and approvals remain mandatory.
6. **Reasoning-plane observation (partially addressed).** The extrapolator reads session transcripts via platform APIs, providing independent observation of tool calls and arguments. However, this is transcript-only – chain-of-thought reasoning, internal model state, and unlogged deliberation are not observable. A sufficiently sophisticated attacker could craft tool calls that appear benign in the transcript while achieving malicious outcomes through side effects invisible to the session record.
7. **Observer integrity boundary.** Section 4.3 assumes a trustworthy observer. Under kernel-level compromise, telemetry can be suppressed or forged.
8. **Single-platform benchmark.** BadAgentUse is currently Linux-only (aarch64 Lima VM). Multiplatform support exists in the observer but is not yet fully benchmarked in this protocol.
9. **Single-observer, non-federated focus.** The current EDAMAME runtime already supports multiple reasoning-plane contributors on

one observer-hosted engine (for example OpenClaw plus Cursor slices merged into one correlation window), but multi-host delegation still requires federated observation and cross-host correlation that are not yet implemented.

10. **No formal safety proofs.** Section 5 provides transparent rule properties, not PAC bounds or cryptographic proofs.
11. **High-impact coding-agent scenarios remain under-tested.** We have not yet fully benchmarked destructive git workflows and secret-exposure paths (e.g., accidental commit/exfiltration of repository credentials and CI tokens), where delayed detection may be especially costly.
12. **External validation gap.** Independent replication remains necessary for claims beyond controlled-environment demonstration.

Benchmark-scope note: the aggregate metrics in Section 7 quantify EDAMAME’s downstream divergence/correlation performance under the explicit-slice producer mode and do not directly measure transcript-to-model accuracy of either the OpenClaw `edamame-extrapolator` skill or EDAMAME’s internal raw-session ingest path. End-to-end reasoning-plane publication is validated separately by the integration and E2E paths described in Section 7.7.

9.2 Near-Term Work

- **Comparative evaluation on shared benchmarks:** run BadAgentUse scenarios against AgentSentinel, LlamaFirewall, and other monitors to produce head-to-head DSR/FPR comparisons. Requires adapting either the benchmark harness or the competing systems to a common telemetry interface.
- **Compound multi-stage attacks:** scenarios where a single compromised action chain touches all five signal dimensions (network, Layer-7, posture, LAN, identity), validating cross-dimension divergence detection in a single evaluation.
- **Adaptive evaluation:** extend BadAgentUse with scenarios that adapt their evasion strategy based on the monitor’s known detection capabilities, following AgentSentinel’s adaptive attack framework.
- **Reasoning-plane auditing:** add a lightweight CoT consistency check as a complementary signal. Before accepting an intent declaration, evaluate whether it is consistent with the conversation context and previous declarations.
- **Payload-aware detection:** extend Example Rule A to consider data volume and connection duration as additional signals for LotL detection.
- **MCP-native poisoning scenarios:** extend BadAgentUse with scenario families inspired by MCPTox (tool-description poisoning) and MSB (tool-response injection). The two-plane approach should detect policy-inconsistent egress or file access induced by such attacks even when the MCP call itself appears legitimate, but this requires empirical validation.

- **Observer-agnostic interface:** define an abstraction layer so that alternative telemetry providers (osquery, raw eBPF, Sysmon) can replace EDAMAME Posture, enabling independent replication and reducing the single-vendor dependency.
- **Ablation baselines:** implement a telemetry-only baseline (B0: blacklist + anomaly score only, no intent gating) and an allowlist-only baseline (B1: no critical signals) to quantify the contribution of intent correlation beyond existing observer signals.
- Measure rollback reliability under reversible action protocols.
- Multi-host replication with federated observers.
- Cross-platform benchmark runs on macOS and Windows (EDAMAME Posture supports both platforms; the benchmark harness requires adaptation).

10. Ethics and Safe Use

This system provides defensive runtime monitoring; it does not grant authorization decisions or replace IAM/policy controls. It should be deployed alongside baseline hardening controls (S1 and S2), manual confirmation paths for high-impact actions, and claim discipline tied to measured evidence and protocol scope.

10.1 Telemetry Privacy and Data Sovereignty

EDAMAME Posture follows a **local-first, user-up** architecture: raw telemetry is collected, processed, and stored on the endpoint. Session data, Layer-7 enrichment, and posture snapshots are retained for a rolling 8-hour window and then discarded. In local-only deployments, telemetry stays on the host. When remote LLM adjudication is enabled, however, a structured evidence bundle may be sent to the configured provider for soft-signal classification and alert synthesis. We therefore avoid the stronger claim that no information ever leaves the host in all modes.

The EDAMAME platform (<https://edamame.tech>) is designed so that endpoint security operates in a user/agent-up model where the user or agent owns and controls their own security posture by default; no administrator can look down into endpoints by default. This stands in contrast to traditional enterprise EDR/MDM architectures where a central admin has full visibility into endpoint activity.

In multi-user or shared-host deployments, Layer-7 enrichment may include usernames and command-line arguments, creating a local privacy surface. The MCP endpoint defaults to localhost binding with PSK authentication (Section 10.2), ensuring telemetry is not exposed beyond the local host.

10.2 MCP Endpoint Exposure

The Streamable HTTP MCP endpoint defaults to localhost binding (127.0.0.1:3000). Misconfiguration – binding to 0.0.0.0 or routing through a reverse proxy without authentication – would expose sensitive runtime telemetry to the network. EDAMAME Posture supports pre-shared key (PSK) authentication; we recommend enabling it whenever the endpoint is reachable beyond loopback. The observer process itself is part of the attack surface (Section 4.3): an adversary who compromises the MCP endpoint can suppress alerts or inject false telemetry. Deployment hardening should follow the MCP specification’s transport security guidelines.

10.3 Ethical Considerations (Menlo Report Alignment)

We map our research to the Menlo Report’s four ICT research ethics principles:

- **Respect for Persons.** No human subjects or real-user data are involved. All evaluation runs in an isolated Lima VM with synthetic injections. Local telemetry can still contain user-adjacent metadata such as usernames or command lines, and remote-provider deployments must therefore treat the provider boundary explicitly. The 8-hour local retention window (Section 10.1) limits how long such data remains on disk.
- **Beneficence.** The system is purely defensive. We document blind spots (LotL evasion, Section 4.3) and design limitations to prevent false confidence.
- **Justice.** Benchmark scenarios, decision rules, and evaluation scripts are open-source for independent replication and critique.
- **Respect for Law and Public Interest.** Telemetry collection occurs only on the researcher’s own VM. No third-party systems are probed or scanned. Benchmark injections target localhost or controlled endpoints only.

10.4 AI Tools Disclosure

Generative AI tools (Claude, Cursor AI) assisted with code scaffolding and paper editing. All outputs were reviewed and validated by the human authors. No AI system is listed as an author. Scenario definitions, decision rules, and final architectural claims were reviewed by the human authors. Deployments may configure `openai/gpt-5.1` or compatible `azure-openai-responses/*` providers for optional adjudication. The local deterministic evidence stage always runs first. In soft-signal cases, the configured LLM may influence the final operator-facing verdict and explanation, but CRITICAL deterministic findings remain non-downgradable. The explicit-slice benchmark reported in Section 7 should therefore be read as a downstream correlation snapshot under controlled model publication, not as proof that all deployments are fully deterministic.

11. Conclusion

We present a two-plane runtime security model: correlate agent intent (reasoning plane) with host behavior (system plane), and flag divergence when behavior is unexplained by intent or triggers intent-independent critical signals (for example CVE-inspired vulnerability checks or blacklisted traffic). We instantiate the model in OpenClaw with EDAMAME Posture and evaluate the downstream correlation loop using trace-backed live artifacts.

The current branch snapshot supports the narrower architectural claim: process-attributed telemetry plus observer-owned intent correlation can yield operationally useful early-warning verdicts. It does not yet justify stronger claims about comparative superiority, benign-noise maturity, or end-to-end transcript-publication accuracy. The boundary conditions remain explicit: living-off-the-land behavior and telemetry timing windows are the primary blind spots, and comparative baseline studies remain future work.

Appendix A. Multi-Signal Scenario Sketches

These sketches illustrate how the same two-plane logic extends beyond pure egress checks when richer telemetry is available.

A1. Lateral movement (LAN + traffic). A compromised skill performs local recon and attempts service probing. LAN discovery APIs detect new hosts/open ports while network sessions attribute traffic to the originating process tree. If reasoning-plane intent remains “document summarization,” the mismatch yields divergence.

A2. Credential theft (identity + traffic). A supply-chain compromise accesses credential stores and initiates outbound transfer. Identity/breach signals increase risk context while Layer-7 open-file/process metadata and network telemetry provide direct evidence for correlation.

A3. Compound attack (network + Layer-7 + identity + posture). A staged workflow scans, reads sensitive material, and exfiltrates. No single signal family is sufficient, but cross-signal correlation yields stage-specific verdicts and supports clearer operator triage.

Appendix B. Portability Notes and Future Work

Section 3.4 and Section 3.5 present the model-level portability argument in the main narrative. This appendix keeps that discussion intentionally brief: it documents the OpenClaw-specific correspondence and then scopes portability work that remains future research.

B.1 How the Skill Architecture Differs from Orchestration Frameworks

Most agentic stacks are built on orchestration frameworks: LangGraph (LangChain, 2024), CrewAI (Moura, 2024), and AutoGen (Microsoft, 2023). These frameworks share several characteristics:

- **Graph-based execution.** Workflows compile into DAG/state-machine steps.
- **Cloud-first telemetry.** Traces often flow through external observability services (e.g., LangSmith, AgentOps, Arize Phoenix).
- **Framework coupling.** Guardrails are usually callback/runtime specific.

The OpenClaw skill model is architecturally distinct on every axis:

Property	Orchestration frameworks	OpenClaw skill model
Execution model	Compiled graph / state machine	Conversational Markdown instructions
Runtime state	Framework-managed checkpoints	Agent-controlled memory files
Security boundary	Framework callbacks / middleware	Declarative manifest (<code>clawhub.json</code>)
Telemetry path	Cloud API (LangSmith, AgentOps)	Local-only MCP (localhost, PSK)
Agent coupling	Framework-specific (LangGraph nodes, CrewAI tasks)	Protocol-level (MCP tools)
Dependency weight	Heavy (Python ecosystem, SDK, cloud auth)	Zero-framework (Markdown + MCP HTTP)

In OpenClaw, the skill is Markdown (`SKILL.md`) interpreted at runtime, without compiled workflow graphs or framework SDK dependencies. This keeps the security contract independent of framework internals.

B.2 Skill-to-Architecture Correspondence

The OpenClaw-facing reasoning skill and the EDAMAME-internal loops (divergence plus posture/remediation) interact with three EDAMAME layers, but through different interfaces depending on whether the component runs inside or outside the observer process:

Telemetry layer. In the reference benchmark deployment, EDAMAME Posture (via `flodbadd`) provides continual, process-attributed sessions with Layer-7 enrichment. On workstation-hosted Cursor paths, the same internal runtime is hosted by the EDAMAME Security app. The internal divergence and vulnerability loops access this data directly via

Rust method calls (`get_current_sessions`, `get_anomalous_sessions`, `get_blacklisted_sessions`, `get_lanscan`) within the host runtime; no MCP round-trip is involved. This is a deliberate security property: the telemetry path from capture to verdict is entirely internal, unreachable by external processes (Section 4.4). The extrapolator skill does not read EDAMAME telemetry at all; it reads only OpenClaw session transcripts via platform APIs (`sessions_list`, `sessions_history`) and publishes its behavioral model to EDAMAME through the `upsert_behavioral_model` MCP tool.

Scoring layer. EDAMAME Foundation computes composite scores and compliance mapping (CIS, ISO 27001, SOC 2). The internal engine reads score directly (`get_score`) for posture-delta detection. The optional posture facade can read the same data via MCP (`get_score`, `advisor_get_todos`) while reusing the same native remediation pipeline.

Action layer. EDAMAME’s agentic pipeline (`agentic_execute_action` / `advisor_undo_action`) provides execute-and-rollback with confirmation gates. The internal posture loop owns scheduled remediation, while the optional posture facade can consume this layer via MCP for explicit invocations. Example Rule B intercepts here to enforce intent consistency before actions execute.

The key insight is architectural alignment with security isolation: the OpenClaw-facing reasoning skill interacts with EDAMAME through authenticated MCP endpoints, while divergence and posture/remediation loops run inside the observer and access telemetry directly. This separation means that compromising the reasoning plane cannot tamper with the telemetry-to-verdict path, while MCP remains the stable, portable interface for external consumers.

B.3 Future Work: Extending Beyond OpenClaw-Class Flows

This paper evaluates a specific class of agent behavior: cooperating single-node conversational workflows sharing state through memory files. We do **not** claim implemented portability to graph-orchestrated, multi-node, or multi-agent frameworks in this draft.

Future work priority areas are:

1. **Node-scoped intent declarations.** Move from session-scoped intent to per-node intent so allowlists can change safely across workflow steps.
2. **Workflow-aware correlation.** Attribute observations to the correct node across transitions and retries, rather than only to the currently active step.
3. **Delegation trust chains.** Verify that delegated intents are constrained by parent intent and do not expand permission budgets across agent boundaries.
4. **Monitorable workflow design.** Enforce explicit node intent, minimal per-node tool scope, observable transitions, bounded delegation depth, and immutable compiled graphs.

5. **Observer-agnostic replication.** Standardize interfaces so alternative observers can provide equivalent system-plane telemetry for independent studies.

B.4 Scope Boundary for This Draft

Accordingly, this manuscript does not claim:

- validated LangGraph/CrewAI/AutoGen portability results,
- formal workflow-level safety proofs for multi-node execution, or
- cloud-governance analysis for deployments where telemetry leaves endpoints.

These remain future-work items after a workflow-aware prototype exists and can be evaluated with trace-backed scenarios.

12. Author Disclosure and Competing Interests

Frank Lyonnet and Antoine Clerget have financial relationships with EDAMAME Technologies (founder and board member respectively). Kave Salamatian is a professor at Université Savoie Mont Blanc, France, and has no financial relationship with EDAMAME Technologies. The bounding argument is that the evaluation target is the correlation method over published artifacts, not proprietary internals, and that independent replication is practically enabled by open code, trace-level evidence, and observer replaceability.

Frank Lyonnet is the founder of EDAMAME Technologies, the company that develops EDAMAME Posture. This paper evaluates EDAMAME Posture as the system-plane observer – a core architectural dependency. Antoine Clerget is a board member of EDAMAME Technologies. Kave Salamatian is a professor at Université Savoie Mont Blanc and contributed to the architectural framing and evaluation methodology; he has no competing interest.

To bound this conflict: (a) the observer is treated as a black-box telemetry provider, not the evaluation target; (b) decision rules, scenarios, harness, and scripts are open-source; (c) every verdict is trace-backed with on-disk artifacts; and (d) the core observer repositories are open-source: **flodbadd** (session capture + Layer-7 attribution), **edamame_foundation** (scoring/threat infrastructure), and **threatmodels** (endpoint threat-model definitions). All are published under EDAMAME Technologies and can be audited or replaced by alternative observers. We actively invite independent replication.

References

Attack Surface and Benchmarks

- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, Mario Fritz. “Not what you’ve signed up for: Compro-

mising Real-World LLM-Integrated Applications with Indirect Prompt Injection.” arXiv:2302.12173 (2023). arXiv:2302.12173

- Qiusi Zhan, Zhixiang Liang, Zifan Ying, Daniel Kang. “InjecAgent: Benchmarking Indirect Prompt Injections in Tool-Integrated Large Language Model Agents.” arXiv:2403.02691 (2024). arXiv:2403.02691
- “Exfiltration of personal information from ChatGPT via prompt injection.” arXiv:2406.00199 (2024). arXiv:2406.00199
- Saeid Jamshidi et al. “Securing the Model Context Protocol: Defending LLMs Against Tool Poisoning and Adversarial Attacks.” arXiv:2512.06556 (2025). arXiv:2512.06556
- Luo et al. “MCPTox: A Benchmark for Tool Poisoning Attack on Real-World MCP Servers.” arXiv:2508.14925 (2025). arXiv:2508.14925
- Wang et al. “MCP Security Bench (MSB): Benchmarking Attacks Against Model Context Protocol in LLM Agents.” arXiv:2510.15994 (2025). arXiv:2510.15994
- Zhou et al. “Technical Report: Evaluating Goal Drift in Language Model Agents.” arXiv:2505.02709 (2025). arXiv:2505.02709

Runtime Defense Systems

- “AgentArmor: Enforcing Program Analysis on Agent Runtime Trace to Defend Against Prompt Injection.” arXiv:2508.01249 (2025). arXiv:2508.01249
- “LlamaFirewall: An open source guardrail system for building secure AI agents.” arXiv:2505.03574 (2025). arXiv:2505.03574
- Li et al. “AegisLLM: Scaling Agentic Systems for Self-Reflective Defense in LLM Security.” arXiv:2504.20965 (2025). arXiv:2504.20965
- Wallarm et al. “A2AS: Agentic AI Runtime Security and Self-Defense.” arXiv:2510.13825 (2025). arXiv:2510.13825
- “MI9: An Integrated Runtime Governance Framework for Agentic AI.” arXiv:2508.03858 (2025). arXiv:2508.03858
- “AgentSentinel: An End-to-End and Real-Time Security Defense Framework for Computer-Use Agents.” arXiv:2509.07764 (2025). arXiv:2509.07764
- Zhang et al. “Pro2Guard: Proactive Runtime Enforcement of LLM Agent Safety via Probabilistic Model Checking.” arXiv:2508.00500 (2025). arXiv:2508.00500
- Chen et al. “MCP-Guard: A Multi-Stage Defense-in-Depth Framework for Securing Model Context Protocol in Agentic AI.” arXiv:2508.10991 (2025). arXiv:2508.10991
- “Securing AI Agent Execution.” arXiv:2510.21236 (2025). arXiv

Formal Verification and Provable Safety

- Xu et al. “ShieldAgent: Shielding Agents via Verifiable Safety Policy Reasoning.” arXiv:2503.22738 (2025). arXiv:2503.22738

- Li et al. “VET Your Agent: Towards Host-Independent Autonomy via Verifiable Execution Traces.” arXiv:2512.15892 (2025). arXiv:2512.15892
- Adharsh Kamath, Sishen Zhang, Calvin Xu, Shubham Ugare, Gagandeep Singh, Sasa Misailovic. “Enforcing Temporal Constraints for LLM Agents (Agent-C).” arXiv:2512.23738 (2025). arXiv:2512.23738
- Roham Koohestani. “AgentGuard: Runtime Verification of AI Agents.” arXiv:2509.23864 (2025). arXiv:2509.23864

Chain-of-Thought Monitoring

- Rogers et al. “Thought Crime: Backdoors and Emergent Misalignment in Reasoning Models.” arXiv:2506.13206 (2025). arXiv:2506.13206
- Baker et al. “Chain of Thought Monitorability: A New and Fragile Opportunity for AI Safety.” arXiv:2507.11473 (2025). arXiv:2507.11473
- “When Chain of Thought is Necessary, Language Models Struggle to Evade Monitors.” arXiv:2507.05246 (2025). arXiv:2507.05246
- Goldowsky-Dill et al. “CoT Red-Handed: Stress Testing Chain-of-Thought Monitoring.” arXiv:2505.23575 (2025). arXiv:2505.23575

Agent Orchestration Frameworks

- LangChain / LangGraph. “LangGraph: Build stateful, multi-actor applications with LLMs.” (2024). docs
- Joao Moura. “CrewAI: Framework for orchestrating role-playing, autonomous AI agents.” (2024). repo
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Chi Wang. “AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation.” arXiv:2308.08155 (2023). arXiv:2308.08155

Standards and Specifications

- Edwin B. Wilson. “Probable Inference, the Law of Succession, and Statistical Inference.” JASA (1927). publisher page
- J. Sobel, D. Friedman. “An Introduction to Reflection-Oriented Programming.” (1996).
- Model Context Protocol (MCP) Security Best Practices (rev 2025-06-18): spec
- Model Context Protocol (MCP) Authorization (rev 2025-06-18): spec
- Inspektor Gadget. “ig-mcp-server: MCP server for eBPF-powered observability.” (2026). repo
- AKS Engineering Blog. “Observe Smarter: Leveraging Real-Time Insights via the AKS-MCP Server.” (2025). article

Incident Sources

- EDAMAME Posture (software): repo

- OpenClaw security advisories: [advisories](#)
- OpenClaw issue #9627 (config secret sprawl): [issue](#)
- Julie Bort. “A Meta AI security researcher said an OpenClaw agent ran amok on her inbox.” TechCrunch (2026). [article](#)
- CVE-2026-25253 (NVD): nvd.nist.gov/vuln/detail/CVE-2026-25253
- CVE-2026-24763 (OpenCVE): app.opencve.io/cve/CVE-2026-24763
- SecurityScorecard STRIKE exposure research (Moltbots): [article](#)
- declawed.io exposed instances dashboard: declawed.io
- VirusTotal on weaponized agent skills (ClawHub): [article](#)
- Palo Alto Networks Unit 42 (memory poisoning): [article](#)